CogLean

A machine-checked library of cognitive theory

Jeff Shrager

October 2026

The problem

Cognitive theories are stated in prose, in simulation code, or as equations fitted to data.

- What follows from a theory's *core assumptions*, as distinct from its parameter settings or implementation details, is rarely checked.
- A simulation shows that a model *can* produce a result, with these parameters and these stimuli.
- It doesn't show which assumptions produced it, whether it survives other parameters, or whether a rival theory would produce it just as well.

Three kinds of evidence a theory can give

One run	**Every run**	**Every model satisfying the postulates**
one model, one parameter setting	one model, all its executions	all models the theory allows
the theory *can* do it	that model *always* does it	the postulates *entail* it
simulation	model checking, sweeps	proof

CogLean is built around the third kind, and keeps the other two in the same language.

Why Lean

- **One definition, run and proved.** The same definition of a memory can be *run* on an experiment's stimuli and *proved about* for all stimuli. The simulation and the theory can't drift apart.
- **Checked.** Every proof step is checked by a small trusted kernel.
- **General.** A theorem can quantify over all nets, all parameter values, all stimulus lists.
- **Cumulative.** Like Mathlib for mathematics: definitions made once, theorems imported rather than re-derived.

The guarantee is relative. Lean checks that a theorem follows from the definitions, not that the definitions capture the theory, still less that the theory is true of people.

What is in the library

- **CogLib**: shared constructs (feature codes, exact probability distributions, discrimination nets, sampled retrieval, attention) and theory-neutral phenomena.
- **EPAM** (Feigenbaum): discrimination-net learning. Stimulus generalization, oscillation, retroactive inhibition, learning to criterion, the serial position effect: all proved. A published serial-anticipation run is reproduced (33–34 of 34 responses agree).
- **Strategy choice** (Siegler & Shrager): distributions of associations. Retrieval is used with probability $1 - (1 - R)^n$; and a finding: a more peaked distribution can *raise* overt strategy use.
- **Multi-digit addition**: the two combined, with 18 predictions neither makes alone.

The generality problem

There are a great many specific models: EPAM, ACT-R, SAM, REM, MINERVA 2, TODAM, connectionist networks...

- Formalizing them one at a time doesn't scale.
- A pile of formalized models doesn't say what any two of them *disagree* about.
- And none of them is "true": unlike mathematics, there's no settled foundation to build on.

Laws and families

Laws: small, theory-neutral properties of a memory.

- *Certainty*: a just-studied item is retrieved for certain.
- *Stability*: studying one item doesn't change another's retrieval.
- *Order-freedom*: the order of study doesn't matter.
- *Recency*: the later of two similar items is favoured.
- *Stored-only*: nothing unstudied is retrieved.
- *Determinacy*: retrieval has a single outcome.

Families: mechanisms, each formalized once with its parameters free.

- Discrimination (EPAM)
- Activation (ACT-R declarative memory)
- Strength with a criterion (DOA)
- Global matching (MINERVA 2)

Predictions are proved *from laws*, not from models.

The theory matrix

	Certainty	Stability	Order-free	Recency	Stored-only	Determinate
Discrimination	**has**	**has**	lacks	lacks	**has**	**has**
Activation	**never**	**never**	**never**	**has**	**has**	**never**
Strength	lacks	lacks	**has**	**never**	**has**	lacks
Global matching	lacks	lacks	**has**	**never**	lacks	**has**

Every cell is a theorem: a general proof that the law holds (**has**) or fails (**never**) at every parameter setting, or a counterexample at one setting (lacks). Every law separates some families, and every family has its own row.

What is true here?

In mathematics, a proved theorem is true. Here there are only competing theories. What can a library build on?

Claim	Example	Status
A law entails a prediction	stability $\Rightarrow$ practiced items stay practiced	proved, permanent
A family has a law	discrimination nets have stability	proved, permanent
People have a law	people lack stability in recall	empirical, revisable

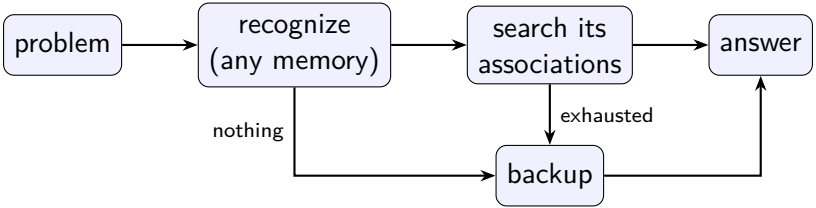
Which theory is right follows from the second and third rows together.

A refutation spreads

- If people lack a law, every family that has it at every setting is refuted at once, and every prediction proved from it loses its support, in every theory that uses it. A counterexample at one setting refutes nothing.
- The matrix carries a judged evidence row, kept apart from the theorems.
- Each judgment states its **bridge**: the assumption connecting what experiments measure (reports, decisions) to what a law is about (what memory retrieves).

Example. People endorse a prototype they never saw. That does not show they *retrieved* it: a memory that returns a similar studied item also calls the prototype "old". The bridge is the missing piece.

Example 1: confident errors, inherited



- **Theorem**: a stranger taken for a familiar item with sharp associations is answered wrongly every time an answer is retrieved, with no more backup use than the familiar item.
- Proved once, for the composition. Inherited by multi-digit addition ($150 + 100$ taken for $100 + 100$) and by face naming (a stranger misnamed).
- From laws: any memory with determinacy, certainty and stability gives it. In the matrix, only discrimination.

Example 2: an experiment derived from laws

Recency and order-freedom disagree. Derive the experiment that separates them.

- Two new items A and B; a probe equally similar to both.
- Study A then B, or B then A; then probe. Order effect Δ = how much more B is favoured after AB than after BA.
- **From laws alone**: order-freedom gives $\Delta = 0$; recency gives $\Delta > 0$; no memory has both.
- With a forced-choice response model, the prediction survives only if guessing is unbiased (or "neither" is equally likely in both orders).

	ACT-R	Strength, MINERVA	EPAM
Δ	10/17 (recency)	0	0, or -2 (primacy)

A related approach: SFT-GMA

Fisher, Houpt, Larue & Schmidt (2025, *J. Math. Psych.*): critical tests of ACT-R's core assumptions, by systems factorial technology (SFT) with global model analysis.

Shared

- Separate a theory's *core* assumptions from auxiliary choices.
- Verdicts per qualitative class, each with the theorem behind it.
- Their finding: ACT-R permits only serial perception within a modality, and data show parallel.

Different

- **Constrains**: processing structure in response times, not what memory retrieves.
- **Data**: SFT has a paradigm, diagnostics and published data; CogLean's link to data is judgment.
- **Scope**: one architecture, not a matrix of families.
- **Checking**: their proofs are careful prose; CogLean's are machine-checked.

Where this is weaker than mathematics

- **Definitions are contestable.** A family is one formalization of a theory. We check that link with blind simulations and community implementations.
- **Laws are idealizations** and presuppose a representation (binary codes, feature-count similarity).
- **Families are flexible.** The informative facts are parameter regions, not single counterexamples.
- **Evidence is never final.** The empirical row will change; the proved rows won't, and the matrix says at once what follows.

Next

- **Learning across trials**: feedback updates memory, so memory states become random; multi-digit learning sequences are the target.
- **Graded laws**, and the parameter regions of the "lacks" cells.
- **More families as tests**: REM, a connectionist memory, CHREST, ACT-R with blending.
- **Observational semantics**: decision and report components, so laws about retrieval predict what participants say.
- **Beyond memory**: strategy choice, then production systems.

Take-home

- Cognitive theories can be stated so that their consequences are *proved*, for every parameter setting.
- The durable content is conditional: laws entail predictions; families have or lack laws.
- A matrix of families against laws tells you which experiments are worth running, and what a result refutes.
- The value is not just that EPAM, ACT-R and others can each be formalized, but that results about one component can be combined with results about another to give predictions none of them anticipated.

github.com/jeffshrager/CogLean